

# CS 5220 Project3 Group 4 Stage 1

Yuan Huang(yh638) Robert Carson(rac428) Ian Vermeulen(iyv2)

For the first stage of the project, we have done the profiling, the MPI parallelization and some tuning on the code.

## 1. Profiling

The first step before attempting to improve performance was to profile the code and identify the hotspots where performance could use the most improvement. To do this, we used VTune Amplifier on the compute nodes. The overall performance measurements are shown below in Figure 1.a.

Function	Module	CPU Time
square	path.x	43.157s
__kmp_barrier	libiomp5.so	12.595s
__kmpc_reduce_nowait	libiomp5.so	4.398s
__kmp_fork_barrier	libiomp5.so	2.844s
__intel_sse3_rep_memcpy	path.x	0.040s
genrand	path.x	0.020s
fletcher16	path.x	0.020s
deinfiniteize	path.x	0.010s
gen_graph	path.x	0.010s

Figure 1.a Overall Profiling of the Code

Unsurprisingly, most of the execution time is spent in the square method. This report also showed us that a significant amount of time was being spent in openMP modules, indicating room for improvement. We drilled down into the square method and the results are shown below in Figure 1.b.

Source Line	Source	CPU Time
37	*	
38	* The return value for `square` is true if `l` and `lnew` are	
39	* identical, and false otherwise.	
40	*/	
41		
42	int square(int n, // Number of nodes	
43	int* restrict l, // Partial distance at step s	
44	int* restrict lnew) // Partial distance at step s+1	
45	{	
46	int done = 1;	
47	#pragma omp parallel for shared(l, lnew) reduction(&& : done)	
48	for (int j = 0; j < n; ++j) {	
49	for (int i = 0; i < n; ++i) {	0.010s
50	int lij = lnew[j*n+i];	0.030s
51	for (int k = 0; k < n; ++k) {	9.279s
52	int lik = l[k*n+i];	26.331s
53	int lkj = l[j*n+k];	
54	if (lik + lkj < lij) {	3.668s
55	lij = lik+lkj;	
56	done = 0;	3.839s
57	}	
58	}	
59	lnew[j*n+i] = lij;	
60	}	
61	}	
62	return done;	
63	}	

Figure 1.b Profiling of square()

## 2. Performance and Speedup

### 2a. MPI implementation

In an attempt to increase the performance of the code, the parallelization model was changed from one that uses a shared memory model (OpenMP) over to a distributed memory model (MPI). In order to use MPI, a 1D domain decomposition was performed along the row indices of the Lnew array. Next since pointers are not valid across processors, the actual contents of the L array need to be passed between processors. Although, when L is being initialized it is desired to only have the master processor initialize the L array. Then the initial values of L can be passed to all the other processors using the MPI\_Bcast call. This means that the contents of Lnew have to be passed to every processors copy of L during the for loop of the shortest path function using the MPI\_Allgatherv command. Overall, the problem with this is that several barriers are placed in the MPI function calls every loop count which will lead to slow downs. Later versions of this code could get around this by using less safe methods of MPI calls that don't rely on the intrinsic use of a barrier. These methods would need to ensure that the data is not touched while it is getting transferred between processors.

### 2b. Performance

The running time of MPI and OpenMP code are shown in the following figure. We can see that the MPI version of code is faster at small size of graph and OpenMP is faster at large size. That's because OpenMP can use shared memory while MPI have to exchange the datas among the processors. When the size of graph gets large, the number of data needed in communication becomes large. So the MPI code spend a lot of time in exchanging the datas among thread. That's why the MPI code works better at small size but performs badly at large size.




Figure 2.b running time vs number of vertices for different code

### 3. Performance models

#### 3a. Weak and Strong Scaling for OpenMP

The original OpenMP version of the code can be analyzed for its strong and weak scaling capabilities. The strong scaling will be tested with a matrix that is 2000x2000. The strong scaling will be defined by Equation 1, and then strong scaling efficiency will be defined by Equation 2. The results can be seen in Table 1.

$$\text{Strong Scaling} = \frac{t_{\text{serial}}}{t_{\text{parallel}}} \quad (1)$$

$$\text{Strong Scaling Efficiency} = \frac{t_{\text{serial}}}{t_{\text{parallel}} p} \quad (2)$$

where  $t_{\text{serial}}$  and  $t_{\text{parallel}}$  are the times taken when running in serial or parallel, respectively, and  $p$  is the number of the threads. The strong scaling efficiency is essentially what percentage of perfect linear scaling is actually achieved. This is plotted on Figure 3.a.1.

Threads	Total Time (seconds)	Strong Scaling	Strong Scaling Efficiency
1	2.95E+01	1.00E+00	100%
2	1.28E+01	2.29E+00	115%
4	6.30E+00	4.68E+00	117%
8	3.55E+00	8.30E+00	104%
12	2.73E+00	1.08E+01	90%
16	2.39E+00	1.23E+01	77%
24	2.89E+00	1.02E+01	43%

Table 1: Strong Scaling for path simulation with 2000 x 2000 matrix size.




Figure 3.a.1: Plot of strong scale efficiency representing data from Table 1.

Based on Figure 3a, it can be seen that after 8 threads that the amount of performance gain obtained by adding more threads goes down. In fact based on Table 1, it can be seen that between 16 and 24 threads that it actually hurts the performance of the code to add more threads.

For weak scaling, five different simulations were performed scaling the number of threads equally with the number of elements per side of the path. The results can be seen in Table 2. Weak scaling will be defined by Equation 3. In these tests, an increase of the number of elements by a factor of  $n$  led to an increase of threads by a factor of  $n^2$ . A plot of the weak scaling can be seen in Figure 3.a.2.

$$Weak\ scaling = \frac{t_{serial}(n(p))}{t_{parallel}(n(p))} \quad (3)$$

Threads	Elements per side	Total time (seconds)	Weak scaling
1	200	1.02E-02	100.0%
2	283	1.95E-02	52.5%
4	400	2.67E-02	38.3%
8	566	5.70E-02	17.9%
16	800	1.53E-01	6.7%

Table 2: Weak scaling for path simulation with varying elements and threads




Figure 3.a.2: Plot of weak scaling representing data from Table 2.

The OpenMP does not currently have terrific weak scaling attributes, and this is due to several different reasons. One of them is that each time the square function is called a new team of threads needs to be created. This is an expensive operation, and the code could be set up such that this only occurs once. Next, a domain decomposition as the ones discussed in class for the water assignment could be of use here as well to better the weak scaling.

### 3b. Weak and Strong Scaling for MPI

The same methodology used for the strong and weak scaling for the OpenMP is used for the MPI codebase. However, it should be noted that since only 12 processors are available on each node a one to one comparison between the OpenMP and the MPI codes. The strong scaling results are given in Table 3. The results of the strong scale efficiency are given in Figure 3.b.1.

Threads	Total Time (seconds)	Strong Scaling	Strong Scaling Efficiency
1	2.70E+01	1.00E+00	100%
2	1.26E+01	2.14E+00	107%
4	9.02E+00	2.99E+00	75%
8	5.26E+00	5.13E+00	64%
10	6.80E+00	3.97E+00	40%
12	4.51E+00	5.98E+00	50%

Table 3: Strong Scaling for path simulation with 2000 x 2000 matrix size.




Figure 3.b.1: Plot of strong scale efficiency representing data from Table 3.

It can be seen when comparing Table 1 and Table 3 that the OpenMP currently has better strong scaling. It can also be seen from Figure 3.b.1 that the strong scale efficiency starts to improve towards the end. One possible reason behind this could be that the MPI framework is better able to communicate across the entire node instead of just portions of it.

The weak scaling part of the codebase still had 5 simulations. Although, it was not able to quite have a one to one comparison with the OpenMP weak scaling. The results are given in Table 4, and then those results are plotted in Figure 3.b.2.

Threads	Elements per side	Total time (seconds)	Weak scaling
---------	-------------------	----------------------	--------------

1	200	8.65E-03	100.0%
2	283	1.91E-02	45.2%
4	400	2.58E-02	33.5%
8	563	5.83E-02	14.8%
12	693	1.11E-01	7.8%

Table 4: Weak scaling for path simulation with varying elements and threads




Figure 3.b.2: Plot of weak scaling representing data from Table 4.

Based on this figure, the weak scaling performs worse than the OpenMP codebase. One reason behind this poor weak scaling could be due to the current number of barriers that the MPI code makes when spreading the update of the L array to all the processors. Then it also has to make a reduce call to determine if all the processors have finished. Improvements could be made by implementing Canon's algorithm into the square function instead of the simple domain decomposition used currently.

## 4.Future works

In the future, we are going to work on the following tunings:

- 1) Try to make the communication in the MPI code non-blocked. We are currently using blocked communications in the code, making the communication unblocked can let us do the calculation work while waiting for the communications to finish.
- 2) Try to do the blocking and copying tunings using the methods in matrix multiplication problem.
- 3) Use Canon's algorithm for the domain decomposition/square part of the code instead of the current simple domain decomposition.